


Exercise 14


You're familiar with the `malloc()/free()` family of functions from the C standard library. These functions get memory from the operating system in big blocks, carve it up to into smaller regions as memory is requested and reorganize memory into large, available blocks as it is freed. The standard library has to manage a data structure to do this, and it has to build that structure out of the same region of memory its using to satisfy memory requests. We're going to try this ourselves.

In the files `myMalloc.h` and `myMalloc.c`, I'm giving you a partial implementation of our own `malloc()`, `free()`, `calloc()` and `realloc()`. Most of this implementation is done for you. It's not nearly as clever as a real heap allocation subsystem would be, but it will be enough to let us see how this kind of thing could work, and why errors in dynamic memory allocation can be particularly tricky to debug. You just have to complete the implementations of two functions, `getBlock()`, which selects and removes a node from the free list, and `returnBlock()` which puts a node back on the free list, potentially merging it with neighboring blocks.

As illustrated below, our implementation keeps free memory organized as a single, big, linked list, linking together all the unused blocks of memory and ordered them by memory address (so, blocks with lower memory addresses will be earlier on the list). The first few bytes of each free block are used as an instance of the **Hole** structure. It contains a next pointer and a field storing the size of this hole. The size measures the number of bytes in the hole, counting the size of the next pointer but not the size of the size field.




When our implementation of `malloc()` needs to allocate a block, it calls the `getBlock()` function. The job of this function is to walk down the linked list of free blocks and unlink and return the first one that's big enough to satisfy the request.



If there's no sufficiently large block (initially, there will be no blocks at all), `getBlock()` should return `NULL`. When `getBlock()` returns `NULL`, our `malloc()` will go to the operating system to ask for more memory. It will use part of that memory to satisfy the request and use `returnBlock()` to put the rest of the memory on the `freeList`.


The block returned by `getBlock()` may be larger than the current request. If there's enough extra space, it may be possible to give away just the first part of the block and re-insert the remaining portion as a new node on the `freeList`.




If the block isn't large enough to make use of the left-over part, we just return the whole block, an example of

internal fragmentation. The `malloc()` function takes care of splitting the block if possible, and it uses `returnBlock()` to put any extra memory back on the free list.


When `malloc()` gives away a pointer to a block of memory, it returns the address right after the size field inside the Hole structure (the orange block in the picture above). The contents of the size field are left in memory, right before the address `malloc()` returns. This is important for when the block is eventually freed. The `free()` function can back up in memory and look in this size field to figure out how much memory is being returned to the free list. It then calls `returnBlock()` to re-link the freed memory back into the right spot in the free list (the block in the figure below is supposed to be a different block from the one allocated above).




The `returnBlock()` function has two jobs. It has to re-link the freed block back into the right spot in the free list, with blocks sorted by memory address. I've already written this part of the function for you.



Then, if there's no gap between the freed block and its predecessor and/or successor blocks on the freeList, `returnBlock()` needs to merge blocks to into a single, larger contiguous block. I've already written part of this code for you. After a block is linked back into the free list, I check to see if there's no gap between the freed block and its successor. If the freed block is adjacent to its successor, I merge them into a single, large block.



The same thing needs to happen with the predecessor. If there's a predecessor block and there's no gap between it and the freed block, they should be merged into a single, large block. You get to add the code to do this.



Be aware that once you start using your own `malloc()/free()`, your program will be using it for **everything**, including memory requests by code you didn't write (e.g., if `printf()` needs to allocate some memory). You may want to debug this code using `gdb`, instead of depending on printouts in your code (since `printf()` may not work if `malloc()` is broken).

To help you try out your substitute heap allocation subsystem, I've written a test program, `memTest.c`. This program uses the `reportFreeList()` function, a debug function that lets us see what's on the free list to make sure it looks correct. This program first tries out some simple allocate/free operations where it's easy to predict what

the free list should look like. Then, it does a series of random malloc()/free() requests to try to find other, less obvious bugs. You should be able to build this test driver with your malloc implementation using a command like:

```
$ gcc -Wall -std=c99 -g -D_BSD_SOURCE memTest.c myMalloc.c -o memTest
```

When you run this test, you should get output like the following. The exact memory address you get for the free list will probably be different, but the sizes will probably be the same. Although all the dynamically allocated memory is freed at the end of randomTest, the free list isn't one giant block like you might expect. It looks like the OS didn't give us all this memory as one big contiguous block, so it shows up as three separate blocks in the free list report.

```
$ ./memTest
----- babyTest 1 -----
----- babyTest 2 -----
0x7f6ec9049000 : 4088
----- toddlerTest 1 -----
0x7f6ec9049080 : 3960
----- toddlerTest 2 -----
0x7f6ec9049028 : 52
0x7f6ec9049080 : 3960
----- toddlerTest 3 -----
0x7f6ec9049000 : 92
0x7f6ec9049080 : 3960
----- toddlerTest 4 -----
0x7f6ec9049000 : 4088
----- randomTest -----
0x7f6ec8eb8000 : 1527800
0x7f6ec9030000 : 98296
0x7f6ec9049000 : 4088
```

Once your replacement malloc is working, submit just the myMalloc.c file for exercise_14. This should be the only file you need to change.